



Cybersecurity

Professor: F. d'Amore

Table of Contents

1 Introduction and terminologies	6
2 Symmetric Encryption	7
2.1 Finite Fields	8
2.2 Per-Round Keys	9
2.3 S-boxes and Bit Shuffles	10
2.4 Feistel Cipher	11
2.5 Padding and Ciphertext Stealing	12
2.6 Stream Ciphers	14
2.7 Types of Stream Ciphers	15
2.8 Stream Ciphers in practice	15
2.8.1 A5/1	16
2.8.2 Rivest Cipher (RC4)	16
2.8.3 Salsa20	16
2.9 Block Ciphers	17
2.10 Block Ciphers in practice	17
2.10.1 Data Encryption Standard (DES)	17
2.10.2 DES Overview	18
2.10.3 DES Complementary Notes	19
2.10.4 Advanced Encryption Standard (AES)	19
2.10.5 AES Complementary Notes	21
2.11 Block Ciphers Modes of Operations	22
2.11.1 Electronic Cook Book (ECB)	22
2.11.2 Cipher Block Chaining (CBC)	23
2.11.3 XEX (XOR Encrypt XOR)	24
2.11.4 XTS (XEX with Ciphertext Stealing)	25
2.11.5 Cipher Feedback (CFB)	26
2.11.6 Output Feedback (OFB)	26
2.11.7 Counter Mode (CTR)	27
3 Data Integrity	29
3.1 Definition	29
3.2 MAC based on CBC Mode Encryption	29
3.3 Vulnerabilities of CBC-MAC	30
3.3.1 Initialization Vector or IV	30
3.3.2 Fixed and Variable-length message	31
3.4 MAC based on cryptographic hash	31
3.5 The Birthday Paradox	32

4 Public Key Cryptography	33
5 Digital Signatures	34
6 Cryptographically Secure Pseudo-Random Number Generators	35
7 Authentication	36
8 Secret Sharing	37
9 Access Control	38
10 Secure Protocols	39
11 Firewalls	40
12 Email Security	41
13 Web Technologies	42
14 Web Security	43
15 Web Tracking	44

List of Figures

1	Feistel Cipher	11
2	CBC CS1 Encrypt	13
3	Basic Structure of DE	18
4	Basic Structure of AE	21
5	Electronic Code Book Encryption	22
6	Electronic Code Book Decryption	23
7	ECB Lack of Diffusion	23
8	Cipher Block Chaining Encryption	24
9	XTS Mode Decryption	25
10	XTS Mode Decryption	26
11	Cipher Feedback (CFB)	26
12	Feedback (OFB)	27
13	Counter Mode (CTR)	28
14	Cipher Block Chaining Residue	30

List of Tables

1 Introduction and terminologies

2 Symmetric Encryption

Secret key cryptography involves the use of a single key. Given a message (the plaintext) and the key, encryption produces unintelligible data which is about the same length as the plaintext was.

Secret key cryptography is sometimes referred to as **conventional cryptography** or **symmetric cryptography**.

Secret key encryption schemes require that both the party that does the encryption and the party that does the decryption share a secret key. We will discuss two types of secret key encryption schemes:

- **Stream Ciphers:** This uses the key as a seed for a pseudorandom number generator, produces a stream of pseudorandom bits, and \oplus s (bitwise exclusive ors) that stream with the data. Since \oplus is its own inverse, the same computation performs both encryption and decryption.
- **Block Ciphers:** This takes as input a secret key and a plaintext block of fixed size (older ciphers used 64-bit blocks, modern ciphers use 128-bit blocks). It produces a ciphertext block the same size as the plaintext block. To encrypt messages larger than the blocksize, the block cipher is used iteratively with algorithms called *modes of operation*. A block cipher also has a decryption operation that does the reverse computation.

So, block ciphers encrypt data in blocks of set lengths, while stream ciphers do not and instead encrypt plaintext one byte at a time. The two encryption approaches, therefore, vary widely in implementation and use cases.

Here we will have some prerequisite concepts and then we will dive into these two encryption algorithms

2.1 Finite Fields

2.2 Per-Round Keys

2.3 S-boxes and Bit Shuffles

Write S-boxes and Bit Shuffles

2.4 Feistel Cipher

It is very important to make the encryption algorithm reversible so that the decryption is possible. One method (used by AES) of having a cipher is to make all components reversible. With DES, the S-boxes are clearly not reversible since they map 6-bit inputs to 4-bit outputs. So instead, DES is designed to be reversible using a clever technique known as a Feistel cipher.

A Feistel cipher builds reversible transformations out of one-way transformations by only working on half the bits of the input value at a time. let us assume a 64-bit input block. (Figure 1) shows both how encryption and decryption work. In a Feistel cipher there is some irreversible component that scrambles the input. We'll call that component the mangler function.

In encryption for round n , the 64-bit input to round n is divided into two 32-bit halves called L_n and R_n . Round n generates as output 32-bit quantities L_{n+1} and R_{n+1} . The concatenation of L_{n+1} and R_{n+1} is the 64-bit output of round n and, if there's another round, the input to round $n+1$. L_{n+1} is simply R_n . To compute R_{n+1} , do the following. R_n and K_n are input to the mangler function, which takes as input 32 bits of data plus some bits of key to produce a 32-bit output. The 32-bit output of the mangler is XORed with L_n to obtain R_{n+1} .

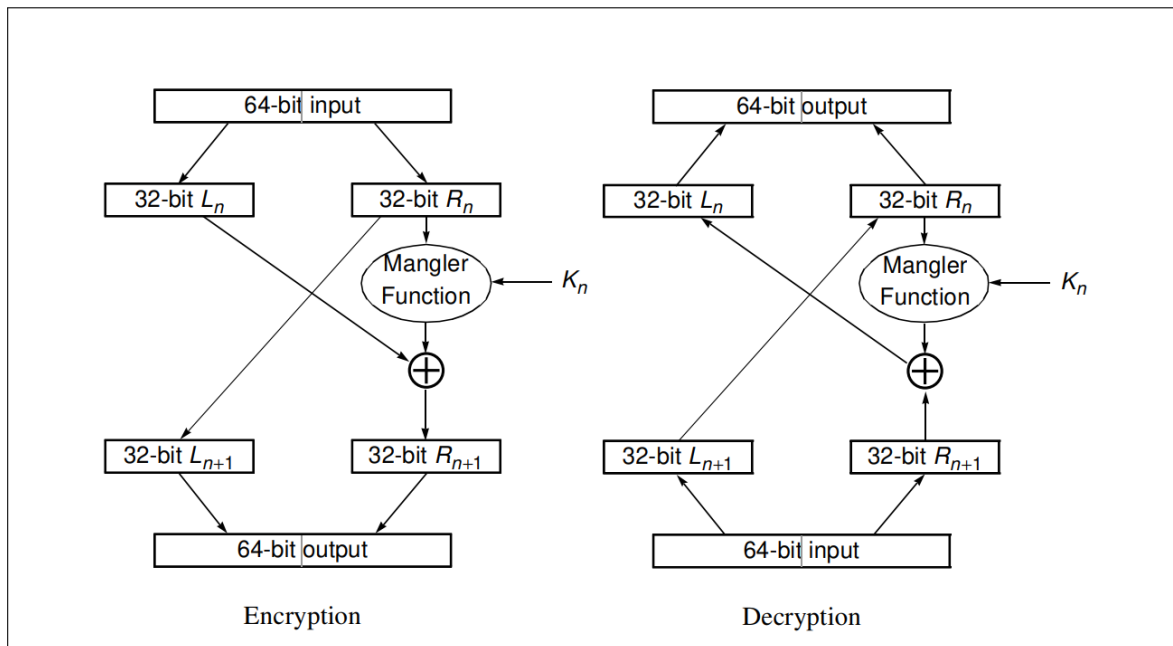


Figure 1: Feistel Cipher

2.5 Padding and Ciphertext Stealing

Padding and ciphertext stealing are techniques used in cryptography to handle messages or plaintext that are not an exact multiple of the block size in block cipher algorithms. Both techniques serve the purpose of ensuring proper encryption and decryption while maintaining the integrity and length of the ciphertext.

Padding:

Padding is the process of adding extra bits or bytes to the plaintext before encryption to ensure it aligns with the block size required by the encryption algorithm. The most commonly used padding scheme is called PKCS#7 (Public Key Cryptography Standard #7).

In PKCS#7 padding, if the plaintext is shorter than the block size, padding bytes are added to the end of the plaintext. The value of each padding byte is set to the number of padding bytes added. For example, if two bytes are needed to reach the block size, both padding bytes will have a value of 0x02.

During decryption, the receiver knows the number of padding bytes based on the value of the last byte in the decrypted block. The padding bytes are then removed to obtain the original plaintext.

Padding is effective when the plaintext length is an exact multiple of the block size, but it introduces additional bytes to the ciphertext, which might impact certain protocols or applications where maintaining the exact plaintext length is crucial.

Ciphertext Stealing:

Ciphertext stealing is an alternative approach to handle the last partial block of plaintext without explicitly padding it. It is commonly used when the plaintext length is not a perfect multiple of the block size.

In ciphertext stealing, the last block of plaintext is divided into two parts: a truncated part and a stolen part.

The truncated part consists of the initial bytes of the last block that can fill up a complete block size. This truncated part is encrypted like any other block and becomes

the second-to-last block of ciphertext.

The stolen part consists of the remaining bytes in the last block. It is then combined with the second-to-last block of ciphertext to create the final block of ciphertext.

During decryption, the last block of ciphertext is decrypted as usual. The stolen part is separated from the decrypted last block, and it is combined with the second-to-last block of plaintext to recover the original message.

Ciphertext stealing allows for encryption and decryption of messages of varying lengths without explicitly padding the last block. It ensures that the ciphertext remains the same length as the original plaintext, which can be desirable in certain scenarios.

Figure (Figure 2) below illustrates the CBC-CS1-Encrypt algorithm for the case that P_n^* is a partial block, i.e., $d \leq b$. The bolded rectangles contain the inputs and outputs. The dotted rectangles provide alternate representations of two blocks in order to illustrate the role of the “stolen” ciphertext. In particular, the string of the $b-d$ rightmost bits of C_{n-1} , denoted P_{n-1}^{**} , becomes the padding for the input block to the final invocation of the block cipher within the execution of CBC mode. The ciphertext that is returned in Step 5 above omits C_{n-1}^{**} , because it can be recovered from C_n during decryption.




Figure 2: CBC CS1 Encrypt

2.6 Stream Ciphers

Idea: try to simulate one-time pad

A stream cipher encrypts a continuous string of binary digits by applying time-varying transformations on plaintext data. Therefore, this type of encryption works bit-by-bit, using keystreams to generate ciphertext for arbitrary lengths of plain text messages. The cipher combines a key (128/256 bits) and a nonce digit (64-128 bits) to produce the **keystream** — a pseudorandom number XORed with the plaintext to produce ciphertext. While the key and the nonce can be reused, the keystream has to be unique for each encryption iteration to ensure security. Stream encryption ciphers achieve this using feedback shift registers to generate a **unique nonce (number used only once) to create the keystream**.

Encryption schemes that use stream ciphers are less likely to propagate system-wide errors since an error in the translation of one bit does not typically affect the entire plaintext block. Stream encryption also occurs in a **linear, continuous manner**, making it simpler and faster to implement. On the other hand, stream ciphers lack diffusion since each plaintext digit is mapped to one ciphertext output. Additionally, they do not validate authenticity, making them vulnerable to insertions. If hackers break the encryption algorithm, they can insert or modify the encrypted message without detection. Stream ciphers are mainly used to encrypt data in applications where the amount of plain text cannot be determined and in low latency use-cases.

2.7 Types of Stream Ciphers

Stream ciphers fall into two categories:

Synchronous Stream Ciphers:

- In a synchronous stream cipher, the keystream block is generated independently of the previous ciphertext and plaintext messages. This means that **each keystream block is generated based only on the key** and does not depend on any previous blocks.
- The most common stream cipher modes use pseudorandom number generators (PRNGs) to create a string of bits, which is combined with the key to form the keystream.
- The keystream is then XORed with the plaintext to generate the ciphertext.

Self-Synchronizing/Asynchronous Stream Ciphers:

- A self-synchronizing stream cipher, also known as ciphertext autokey, **generates the keystream block as a function of both the symmetric key and the fixed-size (N-bits) previous ciphertext block**.
- By altering the ciphertext, the content of the next keystream is changed. This property allows self-synchronizing ciphers to detect active attacks because any modification to the ciphertext will affect the decryption of subsequent blocks, making it easier to detect tampering.
- Asynchronous stream ciphers also provide limited error propagation. If there is a single-digit error in the ciphertext, it can affect at most N bits (the size of the previous ciphertext block) in the next keystream block

.

2.8 Stream Ciphers in practice

popular encryption schemes that use stream ciphers include:

2.8.1 A5/1

2.8.2 Rivest Cipher (RC4)

2.8.3 Salsa20

2.9 Block Ciphers

Block ciphers convert data in plaintext into ciphertext in fixed-size blocks. The block size generally depends on the encryption scheme and is usually in octaves (64-bit or 128-bit blocks). If the plaintext length is not a multiple of 8, the encryption scheme uses **padding** to ensure complete blocks. For instance, to perform 128-bit encryption on a 150-bit plaintext, the encryption scheme provides two blocks, 1 with 128 bits and one with the 22 bits left. 106 Redundant bits are added to the last block to make the entire block equal to the encryption scheme's ciphertext block size.

While Block ciphers use symmetric keys and algorithms to perform data encryption and decryption, they also require an **initialization vector (IV)** to function. An initialization vector is a pseudorandom or random sequence of characters used to encrypt the first block of characters in the plaintext block. The resultant ciphertext for the first block of characters acts as the initialization vector for the subsequent blocks. Therefore, the symmetric cipher produces a unique ciphertext block for each iteration while the IV is transmitted along with the symmetric key and does not require encryption.

Block encryption algorithms offer **high diffusion**; that is, if a single plaintext block were subjected to multiple encryption iterations, it resulted in a unique ciphertext block for each iteration. This makes the encryption scheme relatively tamper-proof since it is difficult for malicious actors to insert symbols into a data block without detection. On the other hand, block ciphers have a **high error propagation rate** since a bit of change in the original plaintext results in entirely different ciphertext blocks.

2.10 Block Ciphers in practice

2.10.1 Data Encryption Standard (DES)

DES is a symmetric block cipher using 64 bit blocks and 56 bits key.

The choice of using 56 bits for keys in DES was a compromise between security

and practicality. While a longer key would provide stronger security, the designers of DES also needed to consider the limitations of computing technology at the time. The inclusion of 8 parity bits reduced the effective key size to 56 bits, which some experts even then considered inadequate for robust security. However, the decision to use a 56-bit key was likely influenced by a balance between security requirements and the feasibility of implementing and processing longer keys with the available hardware during that era.




Figure 3: Basic Structure of DE

2.10.2 DES Overview

First, the 64-bit input is subjected to an initial bit-shuffle. Then, 56-bit key is expanded into sixteen 48-bit per-round keys by taking a 48-bit subset of the 56-bit

input key for each of the keys. Each round takes a 64-bit input with a 48-bit key and produces a 64-bit cipher block. After the sixteenth round, the output has its halves swapped and then is subjected to another bit-shuffle which happens to be the inverse of the initial bit-shuffle. This swapping after the final round does not add any cryptographic strength but has a side benefit. As noted in Feistel Ciphers, Swapping the output make the encryption and description identical except for the key schedule

2.10.3 DES Complementary Notes

The Mangler Function

Undesirable Symmetries

DES Variants

Add DES

Mangler

Function

Add DES

Undesirab

Add DES

Variantss

2.10.4 Advanced Encryption Standard (AES)

AES is a symmetric block cipher using 128 bit blocks and 128, 192 or 256 bits key with the resulting variants called AES-128, AES-192, and AES-256.

AES is similar to DES in that there is a key expansion algorithm which takes the key as an input and expands it into a bunch of round keys, and the algorithm executes a series of rounds that mangle a plaintext block into a ciphertext block. (Figure 4). With DES each round it takes a 64-bit input with a 48-bit key and outputs 64-bit that (along with the next round key) is fed to the next round. With AES, each round takes a 128-bit input and a 128-bit key and produces a 128-bit output which is the input to the next round. Unlike DES, AES is not Feistel cipher. (The Feistel cipher structure is characterized by the division of the plaintext into two halves and the repeated application of a round function that involves both halves. Read more here: 2.4)

In accordance to do as many rounds as needed to make the exhaustive search cheaper than any other form of cryptanalysis, AES does more round with bigger keys. AES-128 has 10 rounds, AES-196 has 12 rounds and AES-256 has 14 rounds. If AES had a 64-bit version, it would have eight rounds, which would be comparable to the sixteen rounds in DES.

Another difference between DES and AES is that DES operates on bits where AES

operates on octets (bytes)

In AES, the key (128, 192, or 256 bits) is expanded into a series of 128-bit round keys, where there is one more round key than there are rounds. AES encryption or decryption consists of \oplus ing a round key into the intermediate value at the beginning of the process, at the end of the process, and between each pair of rounds.

In terms of Substitution and Permutation and cryptographic transformations, DES relies on predefined tables for S-boxes and P-boxes, which may appear arbitrary without understanding their rationale. In contrast, AES provides simpler mathematical formulas for substitutions and permutations, openly stating the reasons behind their design choices.




Figure 4: Basic Structure of AE

2.10.5 AES Complementary Notes

(AES's S-box)

Key expansion

Add AES
Box

Add AES
Key expansion

2.11 Block Ciphers Modes of Operations

Block ciphers operate on fixed-length blocks, typically 64 or 128 bits. The reason for using fixed-length blocks is twofold: first, messages can have varying lengths, and second, encrypting the same plaintext with the same key should always produce the same output.

To address the need for encrypting messages of **arbitrary** lengths, several modes of operation have been invented. These modes allow block ciphers to provide confidentiality for messages of any length. These modes define how the block cipher is applied to multiple blocks and how they are combined to encrypt or decrypt the entire message.

2.11.1 Electronic Cook Book (ECB)

a full 128 bits), and encrypt each block with the secret key (Figure 5). The other side receives the encrypted blocks and decrypts each block in turn to recover the original message (Figure 6).




Figure 5: Electronic Code Book Encryption

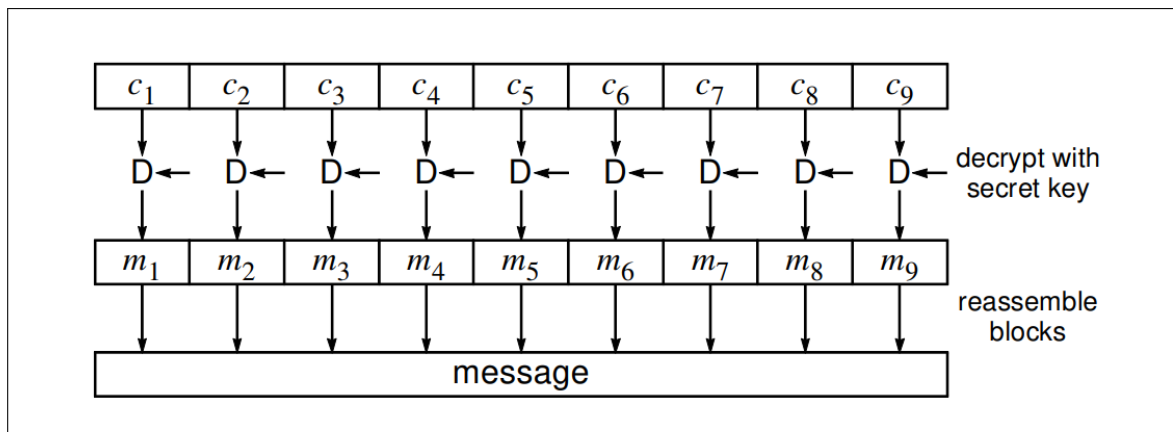


Figure 6: Electronic Code Book Decryption

ECB has two serious flaws. Patterns in the ciphertext, such as repeated blocks, leak information, and nothing prevents someone from rearranging, deleting, modifying, or duplicating blocks.




Figure 7: ECB Lack of Diffusion

2.11.2 Cipher Block Chaining (CBC)

CBC generates its own “random numbers” for all but the first block. It uses c_i as r_{i+1} . In other words, it takes the previous block of ciphertext and uses that as the “random number” that will be \oplus 'd into the next plaintext block. To avoid leaking the information that two messages encrypted with the same key have the same first plaintext blocks, CBC selects one random number that gets \oplus 'd into the first block of plaintext and transmits it along with the data.

This initial random number is known as an **IV (initialization vector)**. A randomly chosen IV guarantees that even if the same message is sent repeatedly, the ciphertext

will be completely different each time.




Figure 8: Cipher Block Chaining Encryption

Properties of CBC:

- What would happen if they changed a block of the ciphertext, say, the value of ciphertext block c_n ? Changing c_n has a predictable effect on m_{n+1} because c_n gets \oplus 'd with the decrypted c_{n+1} to yield m_{n+1} . For instance, changing bit 3 of c_n changes bit 3 of m_{n+1} . Modifying c_n also garbles block m_n to some unpredictable value.
- Asynchronous stream cipher
- Conceals plaintext patterns
- Plaintext cannot be easily manipulated
- No parallel implementation known for encryption
- message must be padded to a multiple of the cipher block size (or we may use Ciphertext Stealing)
- a plaintext can be recovered from just two adjacent blocks of ciphertext
as a consequence, decryption can be parallelized. usually a message is encrypted once, but decrypted many times

2.11.3 XEX (XOR Encrypt XOR)

Write down
XEX

2.11.4 XTS (XEX with Ciphertext Stealing)

XTS is a clever variant of XEX that encrypts multiblock messages that need not be a multiple of the cryptographic blocksize while still keeping the ciphertext the same size as the plaintext. XTS accomplishes this goal (of being length-preserving despite a message not being a multiple of the cryptographic blocksize) through a very clever but somewhat complicated trick known as ciphertext stealing.

Ciphertext stealing pads the final block m_n with as many of the bits of the previous block of ciphertext (c_{n-1}) as necessary to make block m_n be full sized (in our example, $128 - 93 = 35$ bits need to be stolen from c_{n-1}). The padded block m_n is then encrypted. But the ciphertext is now longer than the plaintext, since the last block of ciphertext is now a full-sized block. To solve that problem, we swap ciphertext blocks c_n and c_{n-1} and truncate the final ciphertext block (which was the encryption of block m_{n-1}) to be the size of the original block m_n . In our example, where the original block m_n was 93 bits long, the final ciphertext block will be 93 bits. Now the ciphertext is the same size as the message, but how do we decrypt either of the last two blocks? To obtain block m_n , decrypt c_{n-1} (using implicit IVs and Bmods associated with block m_n). The result will be plaintext m_n with appended padding, but you need to know how much of the result is message and how much is padding. The size of plaintext block m_n will be the size of the final ciphertext block (in our example, 93 bits). The padding (the remaining $128-93=35$ bits) is stolen ciphertext from c_{n-1} .

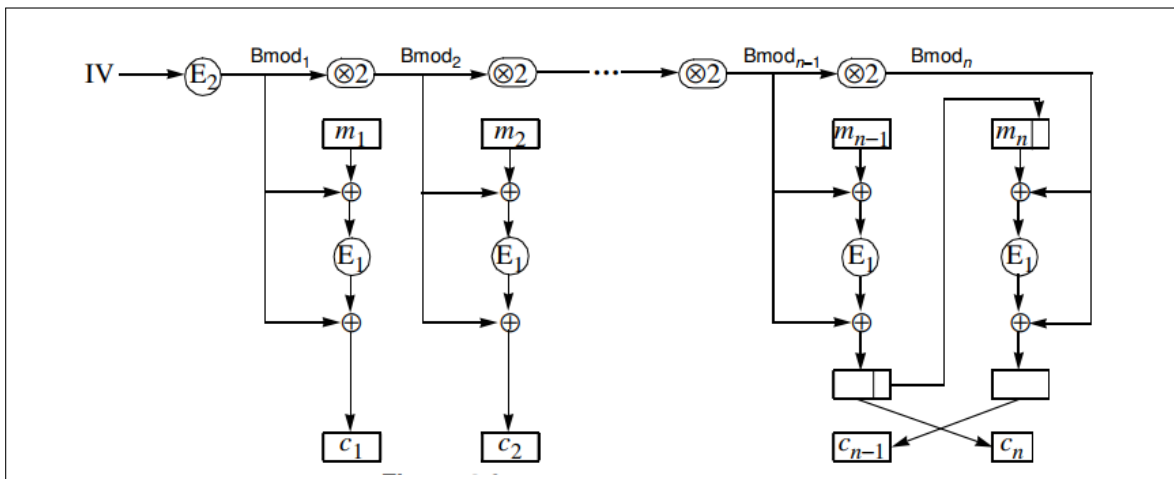


Figure 9: XTS Mode Decryption




Figure 10: XTS Mode Decryption

2.11.5 Cipher Feedback (CFB)

CFB (short for cipher feedback) is an AES block cipher mode like CBC, makes a block cipher into an asynchronous stream cipher (i.e., supports some re-synchronizing after error, if input to encryptor is given through a shift-register) in the sense that for the encryption of a block, B_i , the cipher of the previous block, C_{i-1} is required. CFB also makes use of an initialization vector like CBC. The main difference is that in CFB, the ciphertext block of the previous block is encrypted first and then XOR-ed with the block in focus.

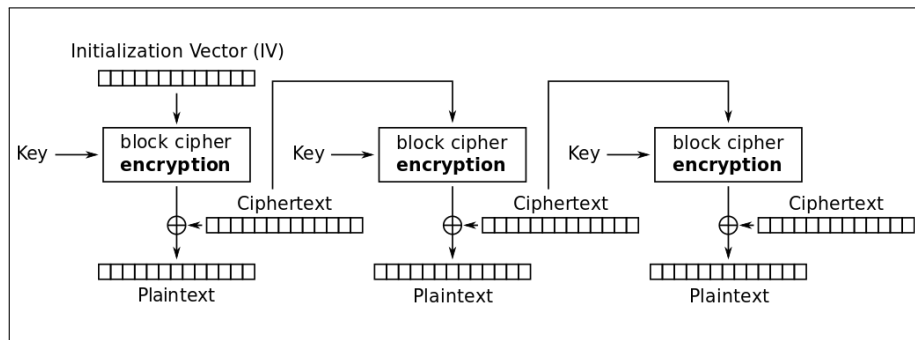


Figure 11: Cipher Feedback (CFB)

2.11.6 Output Feedback (OFB)

The output feedback (OFB) mode makes a block cipher into a synchronous stream cipher. It generates keystream blocks, which are then XORed with the plaintext blocks to get the ciphertext. Just as with other stream ciphers, flipping a bit in the ciphertext produces a flipped bit in the plaintext at the same location. This property

allows many error-correcting codes to function normally even when applied before encryption.

Properties of OFB:

- Synchronous stream cipher
- Errors in ciphertext do not propagate
- Pre-processing is possible
- Conceals plaintext patterns
- No parallel implementation known
- Active attacks by manipulating plaintext are possible

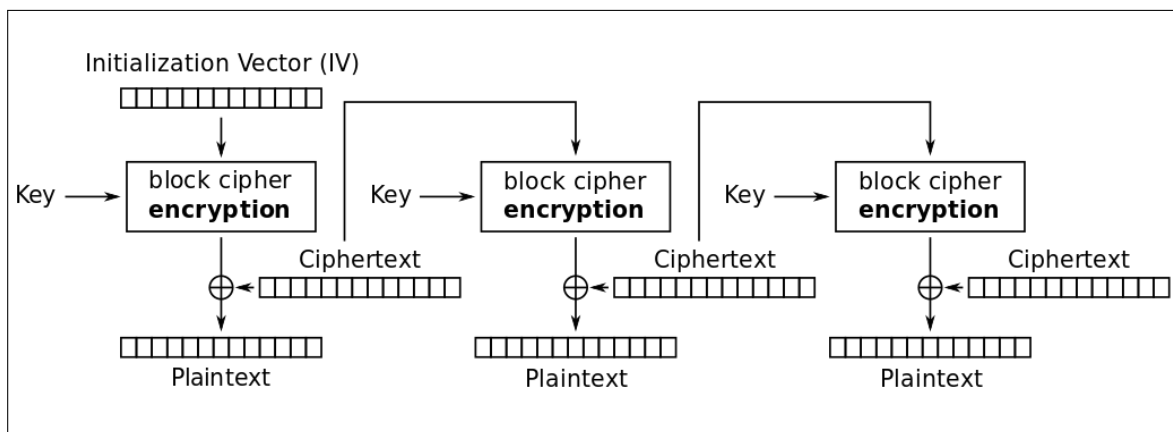


Figure 12: Feedback (OFB)

2.11.7 Counter Mode (CTR)

also known as Integer Counter Mode (ICM) and Segmented Integer Counter (SIC) mode

- turns a block cipher into a stream cipher: it generates the next keystream block by encrypting successive values of a "counter"
- counter can be any function which produces a sequence which is guaranteed not to repeat for a long time, although an actual counter is the simplest and most popular.

- the usage of a simple deterministic input function raised controversial discussions
- has similar characteristics to OFB, but also allows a random access property during decryption
- well suited to operation on a multi-processor machine where blocks can be encrypted in parallel

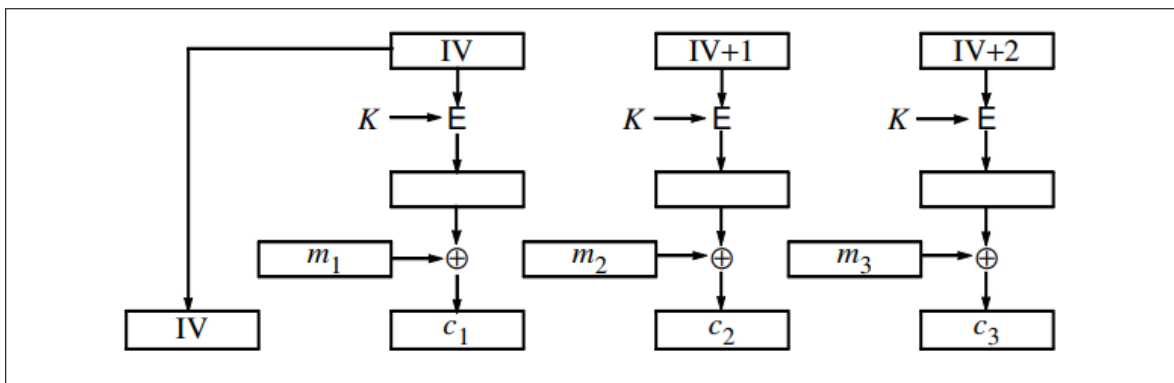


Figure 13: Counter Mode (CTR)

3 Data Integrity

A secret key system can be used to generate a cryptographic integrity check known as a MAC (Message Authentication Code). MACs are used to protect the data from modification in transit.

In this chapter, we'll focus on MACs based on secret key encryption functions and MACs based on hash functions.

3.1 Definition

- *Authentication Algorithm (A)*: A is the algorithm used to generate an authentication tag for a given message using a secret authentication key (k). It takes the message (m) as input and produces the authentication tag $A_k(m)$.
- *Verification Algorithm (V)*: V is the algorithm used to verify the authenticity and integrity of a message. It takes the received message (m) and its corresponding authentication tag ($A_k(m)$) as input and outputs either *accept* or *reject* based on whether the authentication tag is valid or not.
- *Authentication Key (k)*: The authentication key (k) is a secret key shared between Alice and Bob. It is used by the authentication algorithm (A) to generate the authentication tag and by the verification algorithm (V) to validate the tag.
- *Message Space (usually binary strings)*: The message space refers to the set of possible messages that can be authenticated using CBC-MAC. Typically, these messages are represented as binary strings.
- *Message Format*: Every message exchanged between Alice and Bob is a pair $(m, A_k(m))$, where m is the actual message and $A_k(m)$ is its corresponding authentication tag generated by applying the authentication algorithm (A) with the authentication key (k).

3.2 MAC based on CBC Mode Encryption

CBC-MAC computes a MAC of a message using key K . It encrypts the message in CBC mode, using key K , and uses the last block (called the **residue**, see (Figure

14)) as the MAC for the message. If an attacker modifies any portion of a message, the residue will no longer be the correct value (except with probability 1 in 2^{128} , assuming 128-bit blocks).

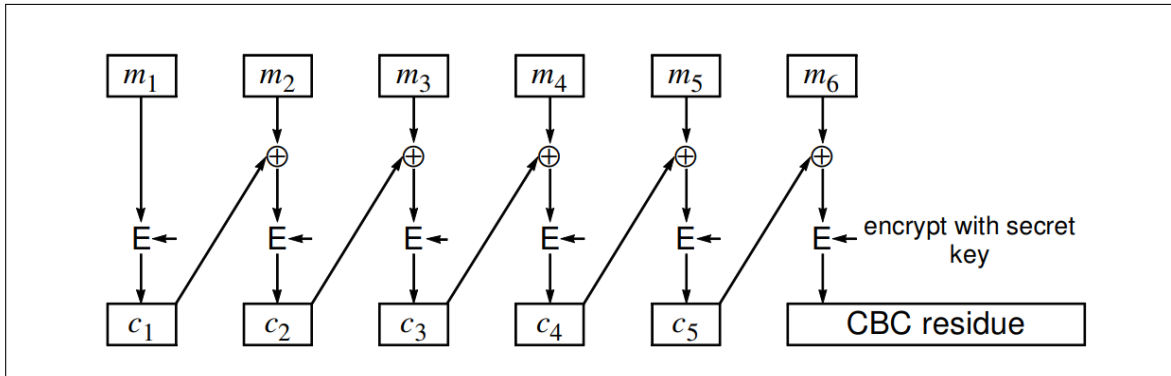


Figure 14: Cipher Block Chaining Residue

3.3 Vulnerabilities of CBC-MAC

3.3.1 Initialization Vector or IV

When we are working on CBC-MAC or Cipher Block Chaining Message Authentication Code, we generally use the Initialization Vector or IV as zero.

When we use the IV or Initialization Vector as zero, a problem arises. Let's say there are two known messages named 'msg1' and 'msg2'. These two messages will generate two signatures called 'sig1' and 'sig2' independently. Finally,

- $E(\text{msg1 XOR } 0) = \text{sig1}$
- $E(\text{msg2 XOR } 0) = \text{sig2}$

Now, a message which consists of msg1 and msg2 will generate two signatures. Let's name the concatenated message as msg3 and the signals generated as sig31 and sig32.

- $E(\text{msg1 XOR } 0) = \text{sig31} = \text{sig1}$
- $E(\text{msg2 XOR sig1}) = \text{sig32}$

We can calculate this without even knowing the key of the encryption.

3.3.2 Fixed and Variable-length message

If the block cipher used is secure (meaning that it is a pseudorandom permutation), then CBC-MAC is secure for fixed-length messages. However, by itself, it is not secure for variable-length messages. Thus, any single key must only be used for messages of a fixed and known length. This is because an attacker who knows the correct authentication tag (i.e. CBC-MAC) pairs for two messages (m, t) and (m', t') can generate a third message m'' whose CBC-MAC will also be t' . This is simply done by XORing the first block of m' with t and then concatenating m with this modified m' ; i.e., by making $m'' = m \parallel (m'_1 \oplus t) \parallel m'_2 \parallel \dots \parallel m'_x$

There are three main ways of modifying CBC-MAC so that it is secure for variable length messages:

1. Input-length key separation
2. Length-prepend
3. Encrypt last block.
4. Use HMACs or CMACs.

3.4 MAC based on cryptographic hash

A hash function inputs an arbitrary-sized bitstring and outputs a fixed-size bitstring, ideally so that all output values are equally likely. A cryptographic hash (also known as a message digest) has some extra security properties:

Preimage Resistance: It should be computationally infeasible to find a message that has a given pre-specified hash.

Collision Resistance: It should be computationally infeasible to find two messages that have the same hash.

Second Preimage Resistance: It should be computationally infeasible to find a second message that has the same hash as a given message.

3.5 The Birthday Paradox

If there are 23 or more people in a room, the odds are better than 50% that two of them will have the same birthday. Analyzing this parlor trick can give us some insight into cryptography.

Let's do this in a slightly more general way. Let's assume n inputs (which would be humans in the birthday example), k possible outputs, and an unpredictable mapping from input to output. With n inputs, there are $\frac{n(n-1)}{2}$ pairs of inputs. For each pair, there's a probability of $\frac{1}{k}$ of both inputs producing the same output value. Therefore, you'll need about $\frac{k}{2}$ pairs in order for the probability to be about $\frac{1}{2}$ that you'll find a matching pair. That means that if n is greater than $\frac{k}{2}$, there's a good chance of finding a matching pair.

4 Public Key Cryptography

5 Digital Signatures

6 Cryptographically Secure Pseudo-Random Number Generators

7 Authentication

8 Secret Sharing

9 Access Control

10 Secure Protocols

11 Firewalls

12 Email Security

13 Web Technologies

14 Web Security

15 Web Tracking